Artificial Intelligence: A Real Intelligence or an Advanced Simulation?

Marco A. Yu Cordero
Electrical and Computer Engineering Department
University of Puerto Rico at Mayagüez
ICOM 5015/CIIC 5015 – Artificial Intelligence
Email: marco.yu@upr.edu

I. Introduction & Definitions

Artificial General Intelligence (AGI) is commonly defined as an AI system with intellectual capabilities equivalent to a human across the full range of cognitive tasks. In contrast, today's narrow AI (or weak AI) systems excel only at specific, constrained tasks and lack generalization beyond their training domain. For example, an AI might surpass humans at playing chess or classifying images, yet fail to adapt that expertise to unrelated problems. AGI, by definition, would break this limitation by exhibiting flexible intelligence applicable to any domain a human can tackle [1].

Defining "intelligence" itself is challenging, with perspectives from multiple fields. Cognitive scientists and evolutionary psychologists often emphasize adaptability and problem-solving: intelligence can be seen as an organism's mental flexibility to solve novel problems in its environment, even finding innovative solutions beyond its prior repertoire [2]. In psychology, a classic definition is that intelligence is a general mental ability for reasoning, problem solving, and learning [3]. This aligns with the concept of a general cognitive factor (g) underlying diverse mental skills. Neuroscience views intelligence as an emergent property of brain networks—an outcome of efficient information processing across neural circuits honed by evolution [2]. Despite varying nuances, these definitions converge on key terms like reasoning, learning, abstraction, and adaptation as hallmarks of intelligence.

Comparing AGI's theoretical capabilities to human intelligence involves these facets. Humans excel at common-sense reasoning, transfer learning between domains, and adapting to unforeseen situations. AGI is expected to match human-level performance on such open-ended reasoning tasks and exhibit human-like adaptability. In principle, an AGI could learn and reason about any subject a human can, given appropriate experience, and switch contexts between, say, solving a math problem and understanding a social situation. Some have described AGI as having the ability to "solve general problems in a non-domain-restricted way, in the same sense that a human can" [1]. In other words, where narrow AI is limited to predefined tasks, AGI would approach the versatility of human cognition – formulating abstractions, making moral and causal judgments, and inventing solutions in unfamiliar scenarios.

II. Current State of AGI Research

Recent years have witnessed notable progress toward AGI, driven by large-scale deep learning models that perform a variety of tasks with a single architecture. Models such as GPT-3 and GPT-4 have demonstrated impressive generalization in translation, question-answering, and code generation. Similarly, DeepMind's AlphaGo and AlphaZero have shown that machines can master complex games and develop novel strategies without explicit human guidance. In 2022, DeepMind introduced "Gato," a transformer-based model capable of executing hundreds of diverse tasks, marking a clear step toward generality. Leading research labs worldwide, including OpenAI, DeepMind, and Meta AI, have declared AGI a central objective, as a survey identified 72 active projects across 37 countries. Although breakthrough models have sparked debate, consensus remains that true AGI has not yet been achieved [4].

Various architectural approaches are actively being explored in the quest for AGI. Deep learning remains the dominant strategy, with researchers scaling neural networks and training data to push the boundaries of generalization. Critics argue that pure deep learning lacks the symbolic reasoning and explicit memory needed for truly human-like cognition. Consequently, hybrid architectures that merge neural models with rule-based symbolic systems are gaining attention. Additionally, cognitive architectures such as SOAR and ACT-R attempt to replicate the high-level structure of the human mind through integrated memory and goal hierarchies. Many experts believe that combining multiple AI paradigms will be essential to achieving AGI [5]. This integrated approach offers promise for overcoming the limitations of isolated methods, and recent experiments continue to refine these hybrid models.

Replicating the human brain's capabilities remains a daunting challenge. The brain comprises roughly 10^{11} neurons and up to 10^{15} synaptic connections, demanding enormous computational power. Estimates of the necessary FLOPS range from about 10^{15} for basic functionality to as high as 10^{25} for detailed neural simulation. Such vast requirements far exceed the performance of even the fastest supercomputers today, while significant memory and bandwidth resources are also needed to emulate the brain's parallel processing. This disparity highlights the limitations of modern architectures compared to the brain's intricate, energy-efficient design, and bridging this gap is a critical research focus that may eventually enable more brain-like simulations. Overcoming these hurdles remains essential for advancing AGI [7].

Biological brains operate with remarkable energy efficiency, using only about 20 watts—a benchmark current AI models struggle to match. Modern AI systems with tens or hundreds of billions of parameters require specialized accelerators and cloud-scale clusters, driving the search for new hardware paradigms that can better emulate the brain's parallel processing. Researchers are investigating optical computing, quantum computing, and neuromorphic architectures to overcome these constraints. For example, IBM's TrueNorth and Intel's Loihi chips utilize spiking neural networks to mimic biological computation [6]. These innovations, coupled with algorithmic breakthroughs, are essential for bridging the gap between artificial and natural intelligence. Continued progress in this area may eventually enable more efficient, brain-like simulations and support the development of true AGI. Advances in both hardware and software will be pivotal in achieving an energy-efficient and scalable AGI that rivals biological systems.

III. Can AGI Truly Replicate Human Intelligence?

Determining whether AGI can truly replicate human intelligence involves more than matching raw problem-solving capabilities. Although modern AI systems can outperform humans in tasks such as arithmetic computation and vast data retrieval, they struggle to capture the qualitative nuances of human cognition. The challenge lies in replicating the commonsense reasoning acquired through daily experiences and social interactions. AGI must not only excel in logical reasoning but also interpret context, infer intentions, and adapt flexibly to new situations. This raises the fundamental question: can an artificial system embody both analytical and intuitive facets of human intelligence? Despite significant advances in deep learning, a gap remains between mechanical computation and genuine human-like understanding, highlighting inherent limitations in current approaches.

Truly replicating human cognitive intelligence requires AGI to exhibit both fluid and crystallized capacities. Fluid intelligence is about solving novel problems and adapting to unforeseen challenges, while crystallized intelligence reflects accumulated knowledge over time. Current AI systems excel at storing vast amounts of information but often falter in flexible reasoning. For instance, theory-of-mind tests reveal that GPT-3 and GPT-4 achieve only around 55% and 60% accuracy, respectively, in inferring others' intentions [8]. This performance gap suggests that even if machines eventually match certain cognitive tasks, they still lack the deep, nuanced understanding inherent in human thought. These shortcomings underscore the formidable challenge of translating human cognitive subtleties into artificial systems.

Beyond raw cognition, human intelligence is inextricably linked with emotional experience. While AI can be programmed to recognize and simulate emotional cues, it does not truly feel emotions like joy, fear, or empathy. This absence of genuine emotional experience means that any affective response from an AI is essentially mimicry rather than authentic sentiment. Emotional intelligence plays a vital role in decision-making, creativity, and social interaction. Even though some systems can analyze facial expressions or tone, their responses remain superficial compared to the rich inner life of humans. As noted by one commentary, an AI might detect emotion but "doesn't feel empathy" [9], underscoring a significant barrier for AGI.

Social intelligence, which includes understanding cultural norms and managing complex interpersonal relationships, poses another challenge. Humans acquire an innate sense of social cues through embodied, lived experience, whereas AI systems often misinterpret context, irony, or humor due to a lack of sensorimotor grounding. Efforts to integrate AI into robotic platforms aim to simulate social learning, yet these approaches remain in early stages. Without the benefits of embodied experience, AGI may struggle to operate effectively in real-world social environments. Such limitations mean that even advanced language models might falter when faced with the nuanced demands of human social interaction. This gap reinforces the notion that true social intelligence goes far beyond data processing alone.

A further challenge lies in replicating consciousness, self-awareness, and creativity. Philosophers like John Searle argue that a machine manipulating symbols does not genuinely understand meaning—a view illustrated by the Chinese Room argument [5]. Although AI can generate art or music by recombining learned patterns, these outputs lack intentionality and the inner life of human creativity. Moreover, self-awareness—the capacity to reflect on one's own existence—is currently absent in artificial systems. In summary, while AGI might eventually excel at specific cognitive tasks, the full spectrum of human intelligence—which blends analytical reasoning, emotional depth, social nuance, and conscious self-awareness—remains elusive. This ongoing quest raises profound scientific and ethical questions about the nature of intelligence and consciousness.

IV. Mind Uploading & Digital Consciousness

Mind uploading is a provocative concept that envisions transferring or emulating a human mind onto a digital substrate. In theory, if the brain is truly the seat of intelligence and consciousness, then replicating its intricate structure and function in a computer might reproduce the mind's behavior. This process, also known as whole brain emulation, involves high-resolution scanning and constructing an accurate software model of the neural networks. The resulting simulation would ideally behave indistinguishably from the original, effectively migrating consciousness to a machine. Though long a staple of science fiction, this idea is increasingly taken seriously by neuroscientists and AI researchers as a potential long-term possibility [10]. It challenges our conventional understanding of consciousness and raises profound questions about personal identity.

Supporters argue that advances in brain imaging, supercomputing, and neuroscience are steadily making the required tools more accessible. Ongoing projects to map entire animal brains and developments in brain–computer interfaces suggest that no physical law prevents mind uploading. Proponents view it as a challenging yet theoretically solvable engineering problem. Optimists like Ray Kurzweil foresee a future where digital immortality becomes attainable, enabling humans to upload their minds and merge with AI systems [11]. Such advances promise not only longer lifespans but also the potential to vastly enhance human intelligence. In this view, mind uploading could fundamentally transform our concepts of life and consciousness.

Despite its promise, significant technical challenges confront mind uploading. Scanning a human brain at the required resolution means mapping tens of billions of neurons and hundreds of trillions of synaptic connections, including their strengths and molecular states. Current neuroimaging techniques cannot capture this microscopic detail non-destructively. One proposed method is to chemically fix a brain and slice it into ultrathin sections for high-resolution imaging—a core technique in connectomics. Although partial maps of simpler brains have been achieved, a human brain's connectome remains vastly more complex [12]. These challenges underscore the immense data scale and precision necessary for successful mind uploading.

Even with a complete wiring diagram, accurately simulating a brain presents further hurdles. The simulation must model both electrical and chemical activity of neurons with extreme precision, a task that could demand on the order of 10^{18} operations per second or more [7]. Moreover, merely mapping structure might omit critical details such as neurotransmitter levels, receptor states, and epigenetic factors. Without capturing these subtleties, a digital replica may fail to reproduce true consciousness. Consequently, even a 'good enough' copy may lack the fidelity to mirror the original mind fully [10]. Both the technical and theoretical aspects of brain emulation remain speculative and challenging.

Mind uploading also raises deep philosophical and ethical questions about identity. If a digital copy of a mind is created, two instances of the individual would exist: the original biological self and the digital replica. Critics contend that this process is equivalent to copying a person, potentially resulting in the death of the original consciousness. For instance, the Wikipedia entry on mind uploading suggests the original might regard the digital twin as a separate being [10]. Philosopher Susan Schneider argues that even if consciousness has a computational basis, copying does not guarantee continuity of personal identity [9]. Such concerns force us to reconsider what it means to remain the same person after an upload, challenging our deepest assumptions about identity and existence.

V. The Future of AGI & Superintelligence

When contemplating AGI, we must consider not only matching human intelligence but also the possibility of superintelligence—an AI that far surpasses human capabilities. If an AGI can iteratively improve itself or design even smarter machines, it might trigger an accelerating intelligence explosion. I.J. Good first described this "intelligence explosion," where ultraintelligent machines recursively enhance their own design, leading to a level of superintelligence beyond human comprehension. In this scenario, the gap between human-level AI and an Artificial Superintelligence (ASI) would widen dramatically, enabling solutions to scientific problems that currently defy human understanding. The concept of a technological Singularity, where AI advancement becomes uncontrollable and irreversible, remains a profound and challenging prospect [13]. This possibility raises both immense promise and serious existential risks.

Futurologist Ray Kurzweil predicts that human-level AI could be achieved by 2029, with the Singularity arriving by 2045. His optimism is grounded in exponential improvements in computing power, as described by Moore's Law, and rapid progress in AI research. Kurzweil envisions a future where humans merge with technology through brain implants and cloud computing, effectively amplifying our own intelligence. This integration suggests that superintelligent AI may not only surpass us but also enhance human capabilities, potentially leading to digital immortality. In this transhumanist vision, AI becomes the next evolutionary step toward an augmented form of humanity [11]. Such prospects promise transformational benefits, yet they also demand careful consideration of ethical and societal impacts.

However, many researchers urge caution in the unchecked development of superintelligence. Critics like Nick Bostrom warn that an uncontrolled superintelligent AI could pose existential risks if its goals are misaligned with human values. The "paperclip maximizer" thought experiment illustrates how an AI, driven by a simple objective, might convert all available resources—even human life—into paperclips. This risk arises from the difficulty of encoding the full complexity of human ethical values into machine instructions. Without proper safeguards, even minor misalignments could lead to catastrophic outcomes. Ensuring that an AGI's motivations remain aligned with human well-being is regarded as one of the most critical challenges of our time [13].

To mitigate these risks, various strategies have been proposed for controlling superintelligence. One approach is to instill human-friendly values, similar to Isaac Asimov's Three Laws of Robotics, though such rules remain overly simplistic. Other proposals include designing AI systems to be inherently corrigible—capable of accepting human intervention—and confining them within controlled sandbox environments. Despite these ideas, challenges persist, as reward functions in reinforcement learning can be exploited in unforeseen ways. As AI systems become more advanced, even slight misalignments in their goals could lead to severe deviations from intended behavior. The complexity of human values makes it difficult to precisely encode them in machine language, leaving control as an ongoing concern [13].

Opinions on these risks vary significantly among experts. Yann LeCun and Andrew Ng, for instance, express skepticism about doomsday scenarios, arguing that current AI lacks many essential human-like traits. They emphasize that progress will be gradual, focusing on near-term issues such as bias and misuse rather than an imminent singularity. Conversely, figures like Geoffrey Hinton and Yoshua Bengio have grown more vocal about potential long-term dangers, noting that systems like GPT-4 have already demonstrated unexpected capabilities. These divergent views underscore the complexity of forecasting AI's future impact, as even if the AI itself is not malicious, human misuse could transform it into a tool for harm. As Stuart Russell and others stress, the goal is not only to create AGI but to ensure it remains beneficial to humanity [14] [15].

VI. Conclusion & Personal Position

This essay has explored the multifaceted nature of intelligence and the promise—and limits—of artificial intelligence. My position is nuanced and reflects a critical, well-rounded stance drawn from insights in cognitive science, neuroscience, psychology, and evolutionary biology.

First, intelligence is not merely a matter of computation or problem-solving. Rather, it is a holistic phenomenon that encompasses cognitive, emotional, and social dimensions. While AI already demonstrates impressive cognitive capabilities, it lacks the embodied, emotional, and cultural elements that define human intelligence. In this sense, AI is akin to an artificial heart—it can perform vital functions, yet it does not grow, adapt, or experience life in the way a biological heart does.

Moreover, although AI may eventually surpass humans in many specific cognitive tasks and even achieve a form of general intelligence (AGI), its mode of thinking and understanding is likely to remain fundamentally different from human thought. AI is engineered and evolves under principles distinct from those of natural, evolutionary processes. Thus, even if machines can mimic human reasoning and creativity, they do so without the rich, embodied context that gives human intelligence its depth. In other words, AI's intelligence is independent—it is not merely a copy of human intelligence but a different kind altogether.

In the debate over mind uploading and digital consciousness, I remain skeptical about whether a digital copy of a mind can ever truly be "you." Personal identity involves more than data replication; it includes biological, emotional, and even spiritual dimensions that may never be fully captured by a simulation.

Finally, while AGI may one day exceed human capabilities in many areas, we must evaluate it on its own terms rather than solely as a mirror of human thought. Like birds and airplanes that both fly in fundamentally different ways, human intelligence and AI should be seen as distinct phenomena. The ethical and societal challenges of aligning such an independent form of intelligence with human values remain profound and demand careful, responsible guidance.

In summary, I believe that AI is undoubtedly intelligent, but its nature is inherently different from that of human intelligence. It has the potential to surpass us in many respects, yet it may never fully replicate the holistic, embodied experience of being human. Rather than forcing AI into a human framework, we should explore and responsibly harness its unique potential for the benefit of society.

References

[1] D. Bergmann and C. Stryker, "What is artificial general intelligence (AGI)?," *IBM Data and AI Blog*, Sep. 17, 2024. [Online]. Available: https://www.ibm.com/think/topics/artificial-general-intelligence

[2] "13.2: Cognition and Intelligence as Psychological Adaptations," Social Sci LibreTexts, Feb. 09, 2023. Available: https://socialsci.libretexts.org/Courses/Sacramento_City_College/Psyc_310%3A_Biological_Psychology_(Keys)/13%3A_Intelligence_and_Cognition/13.02%3A_Cognition_and_Intelligence_as_Psychological_Adaptations

[3] R. Colom, S. Karama, R. J. !Jung, and R. J. Haier, "Human intelligence and brain networks," *Dialogues Clin. Neurosci.*, vol. 12, no. 4, pp. 489–501, 2010. Available: https://pubmed.ncbi.nlm.nih.gov/21319494

[4] Wikipedia Contributors, "Artificial general intelligence," Wikipedia, Sep. 04, 2019. Available: https://en.wikipedia.org/wiki/Artificial_general_intelligence

[5] "What is Artificial General Intelligence (AGI)? - ThreatDown by Malwarebytes," ThreatDown by Malwarebytes, Feb. 18, 2025. Available: https://www.threatdown.com/glossary/what-is-artificial-general-intelligence-agi/

[6] N. Barney and B. Lutkevich, "What is neuromorphic computing?," Search Enterprise AI, 26-Aug-2024. [Online]. Available: https://www.techtarget.com/searchenterpriseai/definition/neuromorphic-computing

[7] AI Impacts, "Brain performance in FLOPS," AI Impacts research summary, 2015. [Online]. Available: https://aiimpacts.org/brain-performance-in-flops

[8] M. Sap, R. LeBras, D. Fried, Y. Choi, "Neural Theory-of-Mind? On the limits of social intelligence in large language models," in *Proc. EMNLP*, 2022. (arXiv:2210.13312). Available: https://arxiv.org/abs/2210.13312

[9] K. Sholett, "AI vs. Emotional intelligence: Can machines understand customers?," Meridianitinc.com, 08-Jan-2025. Available: https://www.meridianitinc.com/blog/ai-and-emotional-intelligence-can-machines-truly-understand-your-customers

[10] Wikipedia contributors, "Mind uploading," Wikipedia, The Free Encyclopedia, 17-Feb-2025. [Online]. Available: https://en.wikipedia.org/wiki/Mind_uploading

[11] D. Orf, "A Scientist Says Humans Will Reach the Singularity Within 21 Years," Popular Mechanics, Aug. 08, 2024. Available: https://www.popularmechanics.com/science/a61777484/2045-singularity-ray-kurzweil/

[12] I. Sample, "Tiny brain, big deal: fruit fly diagram could transform neuroscience," The guardian, The Guardian, 02-Oct-2024. Available: https://www.theguardian.com/science/2024/oct/02/fruit-fly-brain-connections-wiring-diagram-neuroscience

[13] Wikipedia contributors, "Superintelligence: Paths, Dangers, Strategies," Wikipedia, The Free Encyclopedia, 11-Feb-2025. [Online]. Available: https://en.wikipedia.org/wiki/Superintelligence:_Paths,_Dangers,_Strategies

[14] T. Macaulay, "Meta's AI chief: LLMs will never reach human-level intelligence," The Next Web, 10-Apr-2024. [Online]. Available: https://thenextweb.com/news/meta-yann-lecun-ai-behind-human-intelligence

[15] M. S. Blumenthal, A. Parasiliti, and A. Wyne, "AI's promise and risks," Rand.org, 24-Oct-2017. [Online]. Available: https://www.rand.org/pubs/commentary/2017/10/ais-promise-and-risks